

Big Data Processing (ECS76P)

COURSEWORK REPORT



Queen Mary
University of London

Submitted By:

Name: Aman Naredi

Student Id: 230290873

TABLE OF CONTENTS

TASK 1. MERGING DATASETS

1.1 Loading Datasets	2
1.2 Joining Datasets	4
1.3 Converting Date Format	5
1.4 Verifying Datarame Structure	6

TASK 2 AGGREGATION OF DATA

2.1 Counting trips per buisness months	8
2.2 Calculating Platform Profits per Business per Month	12
2.3 Calculating Driver Earnings per Business per Month	15
2.4 Insights	19

TASK 3 TOP K PROCESSING

3.1 Top 5 Pickup Boroughs per Month	20
3.2 Top 5 Dropoff Boroughs per Month	23
3.3 Top 30 Earnest Routes	24
3.4 Insights	26

TASK 4 AVERAGE DATA

4.1 Average Driver Total Pay by Time of Day	27
4.2 Average Trip Length by Time of Day	29
4.3 Average Earning per Mile by Time of Day	30
4.4 Insights	31

TASK 5 FINDING ANOMALIES

5.1 Average Waiting Time per Day in January	32
5.2 Days with High Average Waiting Time	36
5.3 Insights	36

TABLE OF CONTENTS

TASK 6 FILTERING DATA

<i>6.1 Trip Counts with Conditions</i>	37
<i>6.2 Evening Trip Counts by Pickup Borough</i>	38
<i>6.3 Trips from Brooklyn to Staten Island</i>	39

TASK 7 ROUTE ANALYSIS

<i>7.1 Top 10 Popular Routes Analysis</i>	41
---	-----------

TASK 8 GRAPH PROCESSING

<i>8.1 Defining structype of vertex and edge schema</i>	43
<i>8.2 Construct Edges, Vertices DataFrames and Graph</i>	44
<i>8.3 Creating a graph using vertices and edges</i>	46
<i>8.4 Counting connected vertices with the same Borough and same service_zone</i>	47
<i>8.5 Performing page ranking on the graph dataframe</i>	49

Task 1 Merging Datasets

1.1 LOADING DATASETS

Steps and APIs used:

1. *SparkSession Creation:*

- Started the coursework by importing **SparkSession** from **pyspark.sql** and used `SparkSession.builder.appName("NYC").getOrCreate()` to create a **SparkSession** named "NYC" for interacting with Spark.

2. *Defining helper functions:*

- To validate datasets, defined two helper functions: `good_ride_line` and `good_taxi_line`.
 - `good_ride_line`: This function takes a line of text (a row from `rideshare_data.csv`) and validates if it has 15 comma-separated fields. It returns True if valid, False otherwise.
 - `good_taxi_line`: Similar to `good_ride_line`, it validates if a line from `taxi_zone_lookup.csv` has 4 comma-separated fields and returns True if valid.

3. *Accessing S3 data:*

- With reference to the `starterkit.py`, retrieved environment variables for S3 bucket details like `DATA_REPOSITORY_BUCKET`, `S3_ENDPOINT_URL`, `BUCKET_PORT`, `AWS_ACCESS_KEY_ID`, and `AWS_SECRET_ACCESS_KEY`.
- Using this, configured Spark's Hadoop configuration to access the S3 bucket using the retrieved credentials and endpoint details.

4. *Loading rideshare data:*

- Read the "rideshare_data.csv" file from S3 using **`spark.sparkContext.textFile`**.
- **`filter(good_ride_line)`** applies the `good_ride_line` function for filtering.
- The data is transformed into an **RDD** (Resilient Distributed Dataset), which consists of tuples of “,” separated values (rows) using **`map`**.

- The header row, that is the first row of is identified and saved as `header_ride` using **first** and removed from the RDD using **filter**.
 - The RDD is converted into a DataFrame named `ride_df` with column names specified in `column_names` using **toDF**.
5. *Taxi zone data (taxi_zone_lookup.csv):*
- Similar steps are followed as for rideshare data to load "taxi_zone_lookup.csv" from S3 and create a DataFrame named `taxi_df` with column names in `column_names_taxi`.
6. *Cleaning taxi zone data: API's used [**regex_replace**]*
- To clean `taxi_daf`, iterated through each column name in `column_names_taxi`.
 - Used **regexp_replace** for each column to remove double quotes (") from the corresponding column in `taxi_df`.
7. *API's Used:*
- `sparkContext.textFile`
 - `filter`
 - `map`
 - `toDF`.
 - `withColumn`
 - `regexp_replace`

Challenges encountered:

- Removing double quotes from Taxi zone data

Knowledge gained:

- The task showcases techniques for loading data from CSV files, handling headers, transforming data into DataFrames, and exploring their structure.
- Learned cleaning data by removing unwanted characters (double quotes) from a column.

1.2 JOINING DATASETS

Steps and APIs used:

1. *Joining on pickup location:*
 - Used the **join** function from Spark DataFrames to join the `ride_df` and `taxi_df` DataFrames.
 - The join condition is specified as `ride_df.pickup_location == taxi_df.LocationID`. This matches rows in the rideshare data where the pickup location matches the LocationID in the taxi zone data.
 - The resulting DataFrame is named `joined_df_pickup`.
2. *Renaming columns (pickup):*
 - Used a series of **withColumnRenamed** functions to rename the relevant columns obtained from the taxi zone data:
 - "Borough" -> "Pickup_Borough"
 - "Zone" -> "Pickup_Zone"
 - "service_zone" -> "Pickup_service_zone"
 - Also used **drop** ("LocationID") to remove the redundant "LocationID" column from the joined data.
 - The resulting DataFrame is named `joined_df_Renamed`.
3. *Joining on drop-off location:*
 - Performed a second **join** using the same approach as previous step, but this time joining `joined_df_Renamed` (which now includes pickup zone information) with `taxi_df` based on the `dropoff_location` field in the rideshare data.
 - The resulting DataFrame is named `joined_df_dropoff`.
4. *Renaming columns (drop-off):*
 - Similar to step 2, renamed the columns obtained from the taxi zone data for the drop-off location with the prefix "Dropoff_".
 - We again remove the redundant "LocationID" column.
 - The final enriched DataFrame is named `nyc_df`.
 - We use **printSchema()** to display the schema of the final DataFrame.
5. *API'S used:*
 - `join()`
 - `withColumnRenamed()`
 - `drop()`

Knowledge gained:

- Learned how to perform multi-step joins on DataFrames using the join function with appropriate join conditions.
- It showcases techniques for renaming columns and manipulating DataFrame structures.

1.3 CONVERTING DATE FORMAT

Steps and APIs used:

1. *Accessing the date column:*
 - In this task the focus is on the "date" column in the nyc_df DataFrame, accessed date column using **nyc_df.date**
2. *Converting Unix time stamp to date and its data type from string to date, thus creating a new column named date:*
 - Use the **from_unixtime** function from **pyspark.sql.functions** with the **withColumn** function with alias date to create a new column in nyc_df named "date" to convert the UNIX timestamp store in the "date" column.
 - Specified the format string "yyyy-MM-dd" to indicate the desired output format.
 - **date_format** takes the previously converted temporary column and the desired output format string ("yyyy-MM-dd") to ensure the final column has the correct format.
 - Converted the data type of this new "date" column using **to_date** on nyc_df date column from **string** to **date** inside **withColumn** to make changes in the column.
3. *API's used:*
 - from_unixtime
 - withColumn
 - date_format
 - to_date

Challenges encountered:

1. Ensuring the conversion of data type of date column from string to date after converting its format.

Knowledge gained:

- This task demonstrates working with date formats in Spark DataFrames.
- It highlights the use of `from_unixtime`, `date_format` and `to_date` functions for date and its data type conversion.

Output:

```
2024-03-28 16:46:48,770 INFO codegen.CodeGenerator: Code generated in 29.759298 ms
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|business|pickup_location|dropoff_location|trip_length|request_to_pickup|total_ride_time|on_scene_to_pickup|on_scene_to_dropoff|time_of_day|
date|passenger_fare|driver_total_pay|rideshare_profit|hourly_rate|dollars_per_mile|Pickup_Borough|Pickup_Zone|Pickup_service_zone|Dropoff_Bo
rough|Dropoff_Zone|Dropoff_service_zone|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|Uber|25|125|9.8|299.0|2107.0|45.0|2152.0|night|202
3-04-09|51.28|36.02|15.26|60.26|3.68|Brooklyn|Boerum Hill|Boro Zone|Manh
attan|Hudson Sq|Yellow Zone|
|Uber|25|125|3.32|221.0|945.0|121.0|1066.0|morning|202
3-04-09|25.24|13.72|11.52|46.33|4.13|Brooklyn|Boerum Hill|Boro Zone|Manh
attan|Hudson Sq|Yellow Zone|
|Uber|25|125|3.98|321.0|1669.0|47.0|1716.0|afternoon|202
3-04-09|30.23|20.92|9.31|43.89|5.26|Brooklyn|Boerum Hill|Boro Zone|Manh
attan|Hudson Sq|Yellow Zone|
|Uber|25|125|3.68|376.0|2681.0|20.0|2701.0|afternoon|202
3-04-09|33.59|31.07|2.52|41.41|8.44|Brooklyn|Boerum Hill|Boro Zone|Manh
attan|Hudson Sq|Yellow Zone|
|Uber|25|125|3.6|54.0|1081.0|29.0|1110.0|morning|202
3-05-08|22.69|18.89|7.8|61.26|5.25|Brooklyn|Boerum Hill|Boro Zone|Manh
attan|Hudson Sq|Yellow Zone|
|Uber|25|125|3.6|92.0|1003.0|12.0|1015.0|afternoon|202
3-05-08|34.83|18.8|20.67|66.68|5.22|Brooklyn|Boerum Hill|Boro Zone|Manh
attan|Hudson Sq|Yellow Zone|
|Uber|25|125|3.7|1052.0|1363.0|111.0|1474.0|afternoon|202
3-05-08|21.26|20.28|0.98|49.53|5.48|Brooklyn|Boerum Hill|Boro Zone|Manh
attan|Hudson Sq|Yellow Zone|
|Uber|25|125|3.42|431.0|1015.0|16.0|1031.0|night|202
3-05-08|21.42|14.03|7.39|48.99|4.1|Brooklyn|Boerum Hill|Boro Zone|Manh
attan|Hudson Sq|Yellow Zone|
|Uber|25|125|3.87|187.0|1022.0|7.0|1029.0|morning|202
3-05-09|47.48|23.9|23.58|83.62|6.18|Brooklyn|Boerum Hill|Boro Zone|Manh
attan|Hudson Sq|Yellow Zone|
|Uber|25|125|3.8|683.0|1223.0|59.0|1282.0|morning|202
3-05-09|47.22|24.27|29.24|68.15|6.39|Brooklyn|Boerum Hill|Boro Zone|Manh
attan|Hudson Sq|Yellow Zone|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
only showing top 10 rows
```

1.4 VERIFYING DATAFRAME STRUCTURE

Steps and APIs used:

1. *Counting rows:*
 - Used **count** function from Spark DataFrames to get the total number of rows in `nyc_df`.
 - The result is printed to the console using an f-string to include informative text.
2. *Printing schema:*
 - Printed schema (column names and data types) using **printSchema** function of `nyc_df`.
3. *API's used:*
 - `count`
 - `printSchema`

Knowledge gained:

- This task demonstrates using count and printSchema functions to verify DataFrame structure.
- It highlights the importance of checking data quality and consistency after data processing steps.

Output:

Number of rows: 69725864

root

```
-- business: string (nullable = true)
-- pickup_location: string (nullable = true)
-- dropoff_location: string (nullable = true)
-- trip_length: string (nullable = true)
-- request_to_pickup: string (nullable = true)
-- total_ride_time: string (nullable = true)
-- on_scene_to_pickup: string (nullable = true)
-- on_scene_to_dropoff: string (nullable = true)
-- time_of_day: string (nullable = true)
-- date: date (nullable = true)
-- passenger_fare: string (nullable = true)
-- driver_total_pay: string (nullable = true)
-- rideshare_profit: string (nullable = true)
-- hourly_rate: string (nullable = true)
-- dollars_per_mile: string (nullable = true)
-- Pickup_Borough: string (nullable = true)
-- Pickup_Zone: string (nullable = true)
-- Pickup_service_zone: string (nullable = true)
-- Dropoff_Borough: string (nullable = true)
-- Dropoff_Zone: string (nullable = true)
-- Dropoff_service_zone: string (nullable = true)
```

Task 2 Aggregation of Data

2.1 COUNTING TRIPS PER BUSINESS PER MONTH

Steps and APIs used:

1. *Data Loading and Preparation (Covered in Task 1.1 Report)*

- Created `SparkSession`.
- Rideshare data and taxi zone lookup data are loaded from S3 using `textFile` and converted into DataFrames.
- Cleaned `taxi_df` using **`regexp_replace`**.
- DataFrames are joined based on pickup and dropoff locations using `join` and columns are renamed.
- The date column is formatted to "yyyy-MM-dd" format using `date_format`.

2. *Data Aggregation and Visualization:*

- **Converting Data Types:**
 - The date column is converted to a date type using **`to_date`**.
 - `rideshare_profit` and `driver_total_pay` columns are cast to float type using **`cast("float")`**.
- **Extracting Month:**
 - A new column named `month` is added using `month` to represent the month from the date.
- **Counting Trips:**
 - **`groupBy`** is used on `business` and `month` to create groups.
 - **`agg`** with **`count("*").alias("trip_count")`** aggregates the number of trips in each group and assigns an alias (`trip_count`) to the result.
 - `trips_per_business_month` DataFrame stores the aggregated data.

3. *Saving to S3 Bucket:*

- **S3 Resource Object:**
 - Created a resource object for my S3 bucket using the `boto3` library. This object allowed me to interact with S3 for storing and retrieving data.
- **Current Date and Time:**
 - It retrieves the current date and time using **`datetime.now()`** and formats it using **`strftime`**. This timestamp is used to create a unique filename for the output file.
- **Coalescing DataFrame:**

- The **coalesce(1)** function reduces the number of output files written to S3 by merging smaller files into a single file.
- **Output Path:**
 - Defined the S3 path to save the DataFrame. It combines the bucket name, a folder structure (aman_<date_time>), and the filename (trips_per_business_month.csv) with the .csv extension.
- **Saving as CSV:**
 - The **write.csv** method of the DataFrame is used to specify:
 - path: The S3 path constructed earlier.
 - mode: "overwrite" to replace existing files with the same name.
 - header: True to include the column names in the CSV file.

4. Copying to Local Storage:

- Copied trips_per_business_per_month.csv from S3 bucket to my local file system in output folder using command :

```
ccc method bucket cp -r bkt:aman_28-03-2024_17:51:57/trips_per_business_month.csv/ output
```

5. Histogram Generation (Python Script):

- **Data Loading:**
 - pandas is used to read the trips_per_business_month.csv file into a pandas DataFrame named task2_1.
- **Filtering Data:**
 - Separate DataFrames for Uber (task2_1_uber) and Lyft (task2_1_lyft) trips are created by filtering task2_1 based on the business column.
- **Histogram Creation with Matplotlib:**
 - Generated two histograms(barplots) using plt.bar.
 - Each plot focuses on either Uber or Lyft trips with bars representing trip counts for each month.
 - Added titles, labels, and customization (e.g., colors) for clarity.

6. API's used (same in all 3 parts):

- month()
- cast("float")
- groupBy()
- agg()
- count("*")
- sum()

- orderBy()
- coalesce()
- datetime.now()
- strftime("%d-%m-%Y_%H:%M:%S")
- boto3.resource

Challenges encountered:

- Faced problems in saving the files in bucket. Learned about coalesce functionality and merging files.

Knowledge gained:

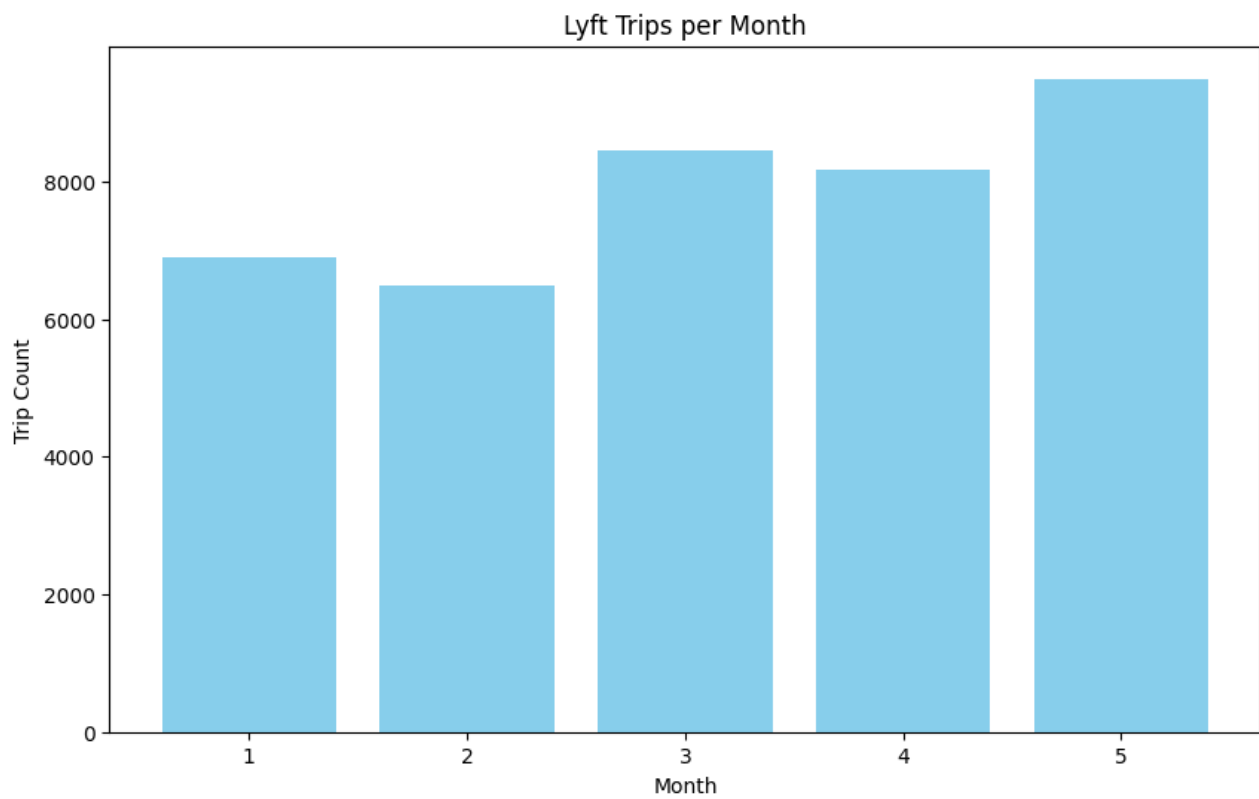

- It highlights the use of aggregation functions like groupBy and count to calculate group-wise statistics.
- Learned the functionality of coalesce API and saving dataframes as csv files into bucket and then how to save them into our local system.

Output:

2024-03-28 17:51:57,123 INFO codegen.CodeGenerator: Code generated in 10.898244 ms

business	month	trip_count
Lyft	4	8173
Uber	5	14276372
Uber	4	13995860
Lyft	3	8444
Uber	2	13280761
Lyft	1	6887
Uber	3	14554308
Uber	1	13579077
Lyft	2	6491
Lyft	5	9491

2024-03-28 17:52:01,481 WARN commit.AbstractS3CommitterFactory: Using standard FileOutputCommitter to commit work. This is slow and potentiall



2.2 CALCULATING PLATFORM PROFITS PER BUSINESS PER MONTH

Steps and APIs Used:

1. Data Loading and Preparation (Done in part 2.1)

2. Calculating Platform Profits:

1. Converting Data Type:

- The `rideshare_profit` column is casted to an float type using `cast("float")` for numerical operations.

2. Profit Aggregation:

- `groupBy` is used on business and month to create groups.
- `agg` with `sum("rideshare_profit").alias("Platform Profit")` calculates the total platform profit for each group (business-month) and assigns an alias (Platform Profit) to the result.
- `trips_per_business_month_profit` DataFrame stores the aggregated data.

3. Saving to S3 Bucket:(similar to previous part)

- **S3 Resource Object:**

- Created a resource object for my S3 bucket using the boto3 library. This object allowed me to interact with S3 for storing and retrieving data.

- **Current Date and Time:**

- It retrieves the current date and time using `datetime.now()` and formats it using `strftime`. This timestamp is used to create a unique filename for the output file.

- **Coalescing DataFrame:**

- The `coalesce(1)` function reduces the number of output files written to S3 by merging smaller files into a single file.

- **Output Path:**

- Defined the S3 path to save the DataFrame. It combines the bucket name, a folder structure (`aman_<date_time>`), and the filename (`trips_per_business_month_profit.csv`) with the `.csv` extension.

- **Saving as CSV:**

- The `write.csv` method of the DataFrame is used to specify:
 - path: The S3 path constructed earlier.
 - mode: "overwrite" to replace existing files with the same name.
 - header: True to include the column names in the CSV file.

4. Copying to Local Storage:

- Copied trips_per_business_per_month.csv from S3 bucket to my local file system in output folder using command :

```
ccc method bucket cp -r bkt:aman_03-04-2024_13:05:20/trips_per_business_month_profit.csv output
```



5. Visualization using Python Script(Similar to previous task):

- *Libraries:*
 - Using **pandas** for data manipulation and **matplotlib.pyplot** for creating histograms.
- *Reading CSV:*
 - Read the downloaded CSV file(trips_per_business_month_profit.csv) into a pandas DataFrame (task2_2).
- *Filtering Data:*
 - Created separate Dataframes for Uber (task2_2_uber) and Lyft (task2_2_lyft) trips by filtering task2_2 based on the business column.
- *Histogram Creation:*
 - The script uses Matplotlib to create separate histograms for Uber and Lyft platform profits, similar to the approach used for trip counts.

Output:

```
2024-04-03 02:12:21,555 INFO scheduler.DAGScheduler: Job 8 finished: snowstring at NativeMethodAccessorImpl.java:0, took 0.099491 s
2024-04-03 02:12:21,593 INFO codegen.CodeGenerator: Code generated in 14.838169 ms
```

business	month	Platform Profit
Lyft	4	-90197.13001759537
Uber	5	1.6313361550055724E8
Uber	4	1.502698201941709E8
Lyft	3	-99403.93998675235
Uber	2	1.3062880563633618E8
Lyft	1	-72633.3500049822
Uber	3	1.520728764191219E8
Uber	1	1.3319711162465689E8
Lyft	2	-70064.72000297531
Lyft	5	-107719.21000343934



2.3 CALCULATING DRIVER EARNINGS PER BUSINESS PER MONTH

Steps and APIs used:

1. *Data Loading and Preparation (Done in part 2.1)*

2. *Calculating Driver Earnings:*

- **Converting Data Type:**
 - The `driver_total_pay` column is casted to an float type using `cast("float")` for numerical operations.
- **Earnings Aggregation:**
 - `groupBy` is used on business and month to create groups.
 - `agg` with `sum("driver_total_pay").alias("Driver Earnings")` calculates the total driver earnings for each group (business-month) and assigns an alias (Driver Earnings) to the result.
 - `trips_per_business_month_driver_pay` DataFrame stores the aggregated data.

3. *Saving to S3 Bucket:(similar to previous part)*

- **S3 Resource Object:**
 - Created a resource object for my S3 bucket using the boto3 library. This object allowed me to interact with S3 for storing and retrieving data.
- **Current Date and Time:**
 - It retrieves the current date and time using `datetime.now()` and formats it using `strftime`. This timestamp is used to create a unique filename for the output file.
- **Coalescing DataFrame:**
 - The `coalesce(1)` function reduces the number of output files written to S3 by merging smaller files into a single file.
- **Output Path:**
 - Defined the S3 path to save the DataFrame. It combines the bucket name, a folder structure (`aman_<date_time>`), and the filename (`trips_per_business_month_profit.csv`) with the `.csv` extension.
- **Saving as CSV:**
 - The `write.csv` method of the DataFrame is used to specify:
 - path: The S3 path constructed earlier.
 - mode: "overwrite" to replace existing files with the same name.
 - header: True to include the column names in the CSV file.

4. Copying to Local Storage:

- Copied trips_per_business_per_month.csv from S3 bucket to my local file system in output folder using command :

```
ccc method bucket cp -r bkt:aman_03-04-2024_15:03:52/trips_per_business_month_driver_pay.csv output
```

5. Visualization using Python Script(Similar to previous task):

- *Libraries:*
 - Using **pandas** for data manipulation and **matplotlib.pyplot** for creating histograms.
- *Reading CSV:*
 - Read the downloaded CSV file(trips_per_business_month_driver_pay.csv) into a pandas DataFrame (task2_2).
- *Filtering Data:*
 - Separate Dataframes for Uber (task2_3_uber) and Lyft (task2_3_lyft) trips are created by filtering task2_3 based on the business column.
- *Histogram Creation:*
 - The script uses Matplotlib to create separate histograms for Uber and Lyft driver earnings, similar to the approach used for trip counts and platform profits.

Challenges encountered:

- *Data type consistency:* Similar to previous tasks, ensuring "driver_total_pay" is an integer type is essential for accurate sum calculation.



Knowledge gained:

- This task reiterates the use of groupBy for group-wise aggregations.
- It highlights the sum function for calculating total values within a group.

Output:

2024-04-04 11:02:58,942 INFO codegen.CodeGenerator: Code generated in 14.635607 ms

business	month	driver_earnings
Lyft	4	297815.3799999999
Uber	5	3.1300511454999596E8
Uber	4	2.9506892721999764E8
Lyft	3	310276.5499999997
Uber	2	2.5215597709000513E8
Lyft	1	239932.2599999999
Uber	3	2.9595849601000154E8
Uber	1	2.5025348066999513E8
Lyft	2	234875.53000000003
Lyft	5	360408.09



2.4 WHEN WE ARE ANALYZING DATA, IT'S NOT JUST ABOUT GETTING RESULTS, BUT ALSO ABOUT EXTRACTING INSIGHTS TO MAKE DECISIONS OR UNDERSTAND THE MARKET. SUPPOSE YOU WERE ONE OF THE STAKEHOLDERS, FOR EXAMPLE, THE DRIVER, CEO OF THE BUSINESS, STOCKBROKER, ETC, WHAT DO YOU FIND FROM THE THREE RESULTS? HOW DO THE FINDINGS HELP YOU MAKE STRATEGIES OR MAKE DECISIONS?

Considering **trip counts** of both the businesses, **Uber** seems to have **stable ridership**, while **Lyft** experiences a **decline**. This presents a considerable **market capture** by **Uber** than **Lyft**.

Uber boasts significant and **increasing platform profits**, while **Lyft** faces consistent **losses**. This highlights Uber's current financial strength compared to Lyft.

Similarly to platform profits, **driver montly earnings** for **Uber** seems to be **increasing** although for **Lyft** it is **fluctuating**. This suggests a need for **Lyft** to consider strategies that **incentivize** driving during peak hours of high-demand periods to retain drivers.

These findings can inform strategic decisions for both companies:

- **Uber:**
 - Analyze Lyft's declining market share and explore targeted campaigns to attract those customers.
 - Continue focusing on strategies that are driving the increase in platform profits.
- **Lyft:**
 - Address the decline in ridership through targeted promotions or adjustments to pricing strategy.
 - Implement cost-cutting measures or explore revenue-generating strategies to address negative profits.

Task 3 Top-K Processing

3.1 TOP 5 PICKUP BOROUGHES PER MONTH

Steps and APIs used:

1. Data Loading and Preparation (Covered in Task 1 Report)

- Created `SparkSession`.
- Rideshare data and taxi zone lookup data are loaded from S3 using `textFile` and converted into `DataFrames`.
- Cleaned `taxi_df` using **`regexp_replace`**.
- `DataFrames` are joined based on pickup and dropoff locations using `join` and columns are renamed.
- The date column is formatted to "yyyy-MM-dd" format using `date_format`.

2. Windowing and ranking:

- Defined a **`window`** specification (`windowSpec`) that partitions data by "month" and sorts them by "trip_count" in descending order.
- Used **`row_number`** function with `windowSpec` to assign a row number to each record within its month, ranking them by trip count (highest gets 1).
- **`filter`** the `DataFrame` to keep only rows with "row_num" less than or equal to 5 (top 5 for each month).
- Finally, **`drop`** the "row_num" column as it's no longer needed.

3. Sorting:

- Used **`orderBy`** to sort the resulting `DataFrame` (`top5_borough_monthly_pickup_trips`) first by "month" (ascending) and then by "trip_count" (descending).
- This ensures the output shows the top 5 boroughs with the highest trip counts for each month.

4. Viewing results :

- Used **`top5_borough_monthly_trips.show(25)`** to display the results.

5. API's used (same in all the parts):

- `window.partitionBy()`: It defines how the data is partitioned before applying a window function. In this scenario, data is partitioned by month, allowing for separate rankings within each month. Used in both the top 5 pickup and dropoff thorough analysis, enabling month-specific ranking.
- `concat_ws()`: Used to either add a new column or replace an existing one. In this task, it's specifically used for adding the Route column and the row_num column post-window function application
- `row_number()`
- `cast("float")`
- `groupBy()`
- `agg()`
- `count("*")`
- `sum()`
- `orderBy()`
- `coalesce()`
- `datetime.now()`
- `strftime("%d-%m-%Y_%H:%M:%S")`
- `boto3.resource`

Challenges encountered:

- Faced challenges in working with window functions and defining other utilities.

Knowledge gained:

- Learned about window functions and using them for ranking data within groups.
- This task highlights techniques for top-k retrievals using row numbering and filtering.

Output:

2024-03-28 19:25:52,197 INFO storage.BlockManagerInfo: Removed broadcast_7_piece0 on 10.133.32.138:37843 in memory (size: 11.0 KiB, free: 2.1 GiB)

Pickup_Borough	month	trip_count
Manhattan	1	5854818
Brooklyn	1	3360373
Queens	1	2589034
Bronx	1	1607789
Staten Island	1	173354
Manhattan	2	5808244
Brooklyn	2	3283003
Queens	2	2447213
Bronx	2	1581889
Staten Island	2	166328
Manhattan	3	6194298
Brooklyn	3	3632776
Queens	3	2757895
Bronx	3	1785166
Staten Island	3	191935
Manhattan	4	6002714
Brooklyn	4	3481220
Queens	4	2666671
Bronx	4	1677435
Staten Island	4	175356
Manhattan	5	5965594
Brooklyn	5	3586009
Queens	5	2826599
Bronx	5	1717137
Staten Island	5	189924

3.2 TOP 5 DROPOFF BOROUGH PER MONTH

Steps and APIs used:

1. *Data Loading and Preparation (Covered in 3.1)*

2. *Window function:*

- Similar to task 3.1, defined a window specification (windowSpec) using the **window** function.
- This window partitions the data by "month" and **sorts** them by "trip_count" in descending order.

3. *Top 5 calculation:*

- Used **row_number** with the defined window to assign a row number to each entry within each month (partition).
- **filter** the DataFrame to keep only rows with a row number less than or equal to 5 (top 5).
- The result is stored in **top5_borough_monthly_dropoff_trips**.

4. *Sorting:*

- Used **orderBy** to sort the final results first by "month" (ascending) and then by "trip_count" (descending).

Output:

```
iB)
2024-03-28 19:47:24,716 INFO storage.BlockManagerInfo: Removed broadcast_9_piece0 on 10.134.100.247:42961 in memory (size: 25.9 KiB, free: 2.1 GiB)
+-----+-----+-----+
|Dropoff_Borough|month|trip_count|
+-----+-----+-----+
|      Manhattan|    1|  5444345|
|      Brooklyn|    1|  3337415|
|       Queens|    1|  2480080|
|       Bronx|    1|  1525137|
|      Unknown|    1|   535610|
|      Manhattan|    2|  5381696|
|      Brooklyn|    2|  3251795|
|       Queens|    2|  2390783|
|       Bronx|    2|  1511014|
|      Unknown|    2|   497525|
|      Manhattan|    3|  5671301|
|      Brooklyn|    3|  3608960|
|       Queens|    3|  2713748|
|       Bronx|    3|  1706802|
|      Unknown|    3|   566798|
|      Manhattan|    4|  5530417|
|      Brooklyn|    4|  3448225|
|       Queens|    4|  2605086|
|       Bronx|    4|  1596505|
|      Unknown|    4|   551857|
|      Manhattan|    5|  5428986|
|      Brooklyn|    5|  3560322|
|       Queens|    5|  2780011|
|       Bronx|    5|  1639180|
|      Unknown|    5|   578549|
+-----+-----+-----+
```

3.3 TOP 30 EARNEST ROUTES

Steps and APIs used:

1. *Creating a route column:*

- Used the **concat_ws** function to create a new column named "Route" that combines the "Pickup_Borough" and "Dropoff_Borough" columns separated by the string "to ".
- This creates a **unique identifier** for each route.

2. *Grouping and aggregation:*

- Used **groupBy** on the "Route" column to group trips for each unique route.
- Calculated the sum of "driver_total_pay" within each group using **sum** to represent the total driver earnings for that route.
- The result is stored in a new DataFrame named **route_profits** with columns "Route" and "total_profit".

3. *Sorting and filtering:*

- Used **orderBy** to sort the **route_profits** DataFrame by the "total_profit" column in descending order (highest profits first).
- Used the **limit** function to select the **top 30** most profitable routes and store the result in **top30_routes**.

4. *Displaying results :*

- Used **show(30, truncate=False)** to display the top 30 routes and their total profits without truncating route names.

Knowledge gained:

- This task demonstrates creating new columns using string manipulation functions.
- It highlights grouping, aggregation, and sorting operations for route-based profitability analysis.

Output:

2024-03-28 19:10:35,977 INFO codegen.CodeGenerator: Code generated in 10.925538 ms

Route	total_profit
Manhattan to Manhattan	3.3385772555001795E8
Brooklyn to Brooklyn	1.739447214800117E8
Queens to Queens	1.1470684719999422E8
Manhattan to Queens	1.0173842820999901E8
Queens to Manhattan	8.603540026000012E7
Manhattan to Unknown	8.010710242000188E7
Bronx to Bronx	7.414622575999315E7
Manhattan to Brooklyn	6.79904755900002E7
Brooklyn to Manhattan	6.3176161049999915E7
Brooklyn to Queens	5.045416242999964E7
Queens to Brooklyn	4.729286535999981E7
Queens to Unknown	4.629299990000313E7
Bronx to Manhattan	3.248632516999994E7
Manhattan to Bronx	3.1978763449999884E7
Manhattan to EWR	2.3750888619999688E7
Brooklyn to Unknown	1.084882756999939E7
Bronx to Unknown	1.0464800209999718E7
Bronx to Queens	1.0292266500000013E7
Queens to Bronx	1.0182898729999995E7
Staten Island to Staten Island	9686862.450000165
Brooklyn to Bronx	5848822.56
Bronx to Brooklyn	5629874.410000006
Brooklyn to EWR	3292761.7099999995
Brooklyn to Staten Island	2417853.819999999
Staten Island to Brooklyn	2265856.4599999976
Manhattan to Staten Island	2223727.3700000001
Staten Island to Manhattan	1612227.72
Queens to EWR	1192758.6600000001
Staten Island to Unknown	891285.8100000016
Queens to Staten Island	865603.3800000001

3.4 SUPPOSE YOU WERE ONE OF THE STAKEHOLDERS, FOR EXAMPLE, EITHER THE DRIVER, CEO OF THE BUSINESS, OR STOCKBROKER, ETC, WHAT DO YOU FIND (I.E., INSIGHTS) FROM THE PREVIOUS THREE RESULTS? HOW DO THE FINDINGS HELP YOU MAKE STRATEGIES OR MAKE DECISIONS?

From above three results it is observable that Manhattan, Brooklyn and Queens are the most famous pickup and drop-off boroughs. Even trips within these boroughs that is Manhattan to Manhattan, Brooklyn to Brooklyn and Queens to Queens are the most profitable trips.

Companies might consider increasing the supply of drivers in these **three** boroughs, to utilize the opportunity of the **high demand**.

The businesses should develop marketing campaigns, offers on fairs or reduced fairs in boroughs like Bronx, Statens Island and EWR. The businesses could expand their services in these boroughs to meet its specific needs, which may have different peak hours or trip purposes.

To ensure an adequate supply of drivers during peak times, the companies could offer incentives or bonuses to drivers who work during the busiest hours or in the most in-demand locations.

Task 4 Average of data

4.1 Average Driver Total Pay by Time of Day

Steps and APIs used:

1. *Data Loading and Preparation (Covered in Task 1 Report)*

- Created `SparkSession`.
- Rideshare data and taxi zone lookup data are loaded from S3 using `textFile` and converted into `DataFrames`.
- Cleaned `taxi_df` using **`regexp_replace`**.
- `DataFrames` are joined based on pickup and dropoff locations using `join` and columns are renamed.
- The date column is formatted to "yyyy-MM-dd" format using `date_format`.

2. *Casting data types:*

- Changed the "trip_length" and "driver_total_pay" columns to float types using **`cast("float")`** to ensure accurate computation for average calculation.

3. *Grouping and aggregation:*

- Used **`groupBy`** on the "time_of_day" column to group trips by their time period.
- Calculated the average of "driver_total_pay" using **`avg`** within each group and stored the result in a new column named **`"average_drive_total_pay"`**.
- The result is stored in **`avg_drive_pay_time_day`** with columns "time_of_day" and "average_drive_total_pay".

4. *Sorting:*

- Used **`orderBy`** to sort **`avg_drive_pay_time_day`** by the "average_drive_total_pay" column in **`descending`** order, showing the time periods with the highest average earnings first.

5. *API's used (same in all the parts):*

- `orderBy ()`

- `agg()`
- `avg()`
- `join()`
- `withColumn()`
- `cast("float")`
- `groupBy()`

Challenges encountered:

- Ensuring consistent data types (e.g., float for numerical values) for aggregation operations like calculating averages.

Knowledge gained:

- Learned changing datatypes of columns using `cast`.
- It highlights using `avg` for calculating average values within groups.

Output:

2024-04-03 13:34:36,123 INFO codegen.CodeGenerator: Code generated in 10.831996 ms

time_of_day	average_drive_total_pay
afternoon	21.212428755696347
night	20.08743800270718
evening	19.777427701749236
morning	19.63333279274821

4.2 AVERAGE TRIP LENGTH BY TIME OF DAY

This task calculates the average trip length for each time of day period to identify the period with the highest average trip distance.

Steps and APIs used:

1. Grouping and aggregation:

- Used **groupBy** on the "time_of_day" column to group trips by their time period.
- Computed the average "trip_length" using **avg** within each group and stored the result in a new column named "average_trip_length".
- The result is stored in **avg_trip_length_time_day** with columns "time_of_day" and "average_trip_length".

2. Sorting:

- Used **orderBy** to sort **avg_trip_length_time_day** by the "average_trip_length" column in **descending** order, showing the time periods with the highest average trip lengths first.

Output:

```
2024-04-03 13:41:15,470 INFO codegen.CodeGenerator: Code generated in 9.733817 ms
2024-04-03 13:41:15,470 INFO codegen.CodeGenerator: Code generated in 9.733817 ms
```

time_of_day	average_trip_length
night	5.323984802300155
morning	4.927371866627282
afternoon	4.861410525884578
evening	4.484750367647451

4.3 AVERAGE EARNING PER MILE BY TIME OF DAY

Steps and APIs used:

1. *Joining DataFrames:*

- Used the **join** function to combine the **avg_drive_pay_time_day** and **avg_trip_length_time_day** DataFrames based on the matching "time_of_day" column.
- This creates a new DataFrame named **joined_df** that includes columns from both DataFrames.

2. *Calculating average earning per mile:*

- Added a new column named "**average_earning_per_mile**" to **joined_df** using **withColumn()** API.
- The value in this column is calculated by **dividing** the "**average_drive_total_pay**" by "**average_trip_length**". This provides the average earnings per unit of distance traveled during each time period.

3. *Selecting and displaying results:*

- Used the **select** function to choose only the "time_of_day" and "average_earning_per_mile" columns for the final result.
- Displayed the results using **show(truncate=False)** to display output without truncating the time-of-day values.

Knowledge gained:

- Learned joining DataFrames based on a shared column.
- It highlights creating new columns with calculations involving existing columns.

Output:

```
2024-04-03 13:52:07,242 INFO scheduler.DAGScheduler: DAG is finished. Submitting to nativeExecutionEnvironmentImpl.java:0, took 0.11700 s
2024-04-03 13:52:07,242 INFO codegen.CodeGenerator: Code generated in 11.431698 ms
```

time_of_day	average_earning_per_mile
afternoon	4.3634308690349854
night	3.7730081412006795
morning	3.9845445653743488
evening	4.409928330553616

4.4 WHAT DO YOU FIND (I.E., INSIGHTS) FROM THE THREE RESULTS? HOW DO THE FINDINGS HELP YOU MAKE STRATEGIES OR MAKE DECISIONS?

From the three results it is observable that:

- The 'afternoon' period recorded the **second highest “average driver total pay”** compared to other times of the day and accounts **second least “average trip length”** thereby resulting into **highest “average earning per mile”**.
- From this we can infer that **afternoon** period is a **high peak fare time** and **more profitable** period for the **company and drivers** as for **short trip length** the **pay is high**.
- Alternatively, **another** reason for **high average driver total pay** and **less average trip length** for **afternoon** period can be the number of rides in **afternoon period** being **less**, might be because of **high fares**. So here, **companies** can adjust the fares to **increase** trips in afternoon.
- The 'night' time has the **“highest average trip length”**. But accounts to **least “average earning per mile”**. Thus this is the **period** where the **companies** should **increase the fares** to make **profit**.

Task 5 Finding anomalies

5.1 AVERAGE WAITING TIME PER DAY IN JANUARY

Steps and APIs used:

1. Data Loading and Preparation (Covered in Task 1 Report)

- Created `SparkSession`.
- Rideshare data and taxi zone lookup data are loaded from S3 using `textFile` and converted into `DataFrames`.
- Cleaned `taxi_df` using `regexp_replace`.
- `DataFrames` are joined based on pickup and dropoff locations using `join` and columns are renamed.
- The date column is formatted to "yyyy-MM-dd" format using `date_format`

2. Filtering data:

- The `filter` function with `month("date") == 1` selects only those rides that occurred in January.
- This creates a new `DataFrame` `jan_df` containing January data.

3. Grouping and aggregation:

- `groupBy` is used on the `dayofmonth("date")` column, which extracts the **day of the month** from the "date" column and assigns it an alias "day".
- Calculated the average "**request_to_pickup**" time using `avg` within each day group and stored the result in a new column named "**average_waiting_time**".
- The result is stored in `avg_waiting_time_per_day`, showing average waiting time for each day in **January**.

4. Sorting and Displaying first 20 rows:

- `orderBy("day")` to sort `avg_waiting_time_per_day` by the "day" column, ensuring results are presented chronologically.
- Printed first 20 rows of resulting dataframe in output using `avg_waiting_time_jan.show(20)`.

5. Saving to S3 Bucket:(similar to Task 2)

- **S3 Resource Object:**
 - Created a resource object for my S3 bucket using the boto3 library. This object allowed me to interact with S3 for storing and retrieving data.
- **Current Date and Time:**
 - It retrieves the current date and time using **datetime.now()** and formats it using **strftime**. This timestamp is used to create a unique filename for the output file.
- **Coalescing DataFrame:**
 - The **coalesce(1)** function reduces the number of output files written to S3 by merging smaller files into a single file.
- **Output Path:**
 - Defined the S3 path to save the DataFrame. It combines the bucket name, a folder structure (aman_<date_time>), and the filename (trips_per_business_month_profit.csv) with the .csv extension.
- **Saving as CSV:**
 - The **write.csv** method of the DataFrame is used to specify:
 - path: The S3 path constructed earlier.
 - mode: "overwrite" to replace existing files with the same name.
 - header: True to include the column names in the CSV file.

6. Copying to Local Storage:

- Copied avg_waiting_time_jan.csv from S3 bucket to my local file system in output folder using command :


```
ccc method bucket cp -r bkt: aman_29-03-2024_02:14:32/avg_waiting_time_jan.csv/ output
```

7. Visualization using Python Script(Similar to previous task):

- **Libraries:**
 - Using **pandas** for data manipulation and **matplotlib.pyplot** for creating histograms.
- **Reading CSV:**
 - Downloaded and read the csv file(avg_waiting_time_jan.csv) into a pandas DataFrame (task5).
- **Histogram Creation:**
 - Generated a Histogram (barplot) for visualising average waiting time of January using Matplotlib library function **plt.bar()** with “day” on **x_axis** and **average waiting time** on “y-axis”.

8. *API's used (same in all parts):*

- filter()
- month()
- groupBy()
- agg()
- avg()
- dayofmonth()
- orderBy()
- filter()
- write.csv()

Challenges encountered:

The challenge of saving the dataframe as a csv file into our system, which was solved while completing task2.

Knowledge gained:


- Learned data filtering methodology based on specific conditions (month).
- It highlights using dayofmonth to extract the day of the month from a date column.

Output:

2024-03-29 02:14:32,502 INFO codegen.CodeGenerator: Code generated in 8.100081 ms

day	average_waiting_time
1	396.5318744409635
2	246.05148716456986
3	235.68026834234155
4	228.85434668408274
5	226.08877381422872
6	230.35306927438575
7	233.25699185710533
8	246.41358687741243
9	229.265944341545
10	225.65276195086662
11	224.40468798627612
12	255.17599322195403
13	239.22308233638282
14	247.49345781069232
15	268.5346481777792
16	251.55102299494047
17	240.5772885527869
18	231.90770494488552
19	272.02203820618143
20	243.43761253646377

only showing top 20 rows



5.2 Days with High Average Waiting Time

Steps and APIs used:

1. *Filtering results:*

- Applied **filter** function on **avg_waiting_time_jan** dataframe (created in Task 5.1) to select only days where the "average_waiting_time" is greater than 300 seconds.
- This filtered DataFrame, named **day_greater_300_secs**, highlights days with potentially high waiting times.

2. *Displaying results:*

- Displayed the content of **day_greater_300_secs**, which includes the result of filter operation, "day" and "average_waiting_time" columns using **day_greater_300_secs.show()** .

Output:

```
2024-03-28 12:00:33,066 INFO codegen.CodeGenerator: Code generated in 10.11413 ms
+---+-----+
|day|average_waiting_time|
+---+-----+
| 1 | 396.5318744409635 |
+---+-----+
```

5.3 Why was the average waiting time longer on these day(s) compared to other days?

The average wait time exceeds 300 only on 1st January , likely due to New Year celebrations. Since many individuals attend parties to ring in the New Year, there is heightened demand on this day, resulting in longer wait times compared to other days of January, where the average waiting time typically hovers around 270. This presents a reasonable anomaly in comparison to the rest of the month.

Task 6 Filtering Data

6.1 TRIP COUNTS WITH CONDITIONS

Steps and APIs used:

1. Data Loading and Preparation (Covered in Task 1 Report)

- Created `SparkSession`.
- Rideshare data and taxi zone lookup data are loaded from S3 using `textFile` and converted into `DataFrames`.
- Cleaned `taxi_df` using `regexp_replace`.
- `DataFrames` are joined based on pickup and dropoff locations using `join` and columns are renamed.
- The date column is formatted to "yyyy-MM-dd" format using `date_format`

2. Grouping and aggregation:

- Used `groupBy` on both "Pickup_Borough" and "time_of_day" columns.
- Count the number of trips using `count` within each group and store the result in a new column named "trip_count".
- This creates `borough_time_of_day_trip_counts`, showing trip counts for each combination of pickup borough and time of day.

3. Filtering:

- Used `filter` with a compound condition to `select` only rows where "trip_count" is greater than 0 but less than 1,000.
- This ensures the focus on trip counts within the specified range.

4. Displaying results:

- Displayed the filtered `DataFrame` `filtered_trip_counts` using `show(truncate=False)` to print without truncating column values, providing a clear view of pickup boroughs, times of day, and their corresponding trip counts.

5. API's used (same in all parts):

- `groupBy()`
- `agg()`

- `filter()`
- `select()`
- `show()`
- `count()`
- `col()`

Knowledge gained:

- Learned filtering DataFrames based on multiple conditions applied to columns.
- Also about grouping by two columns for the required analysis.

Output:

2024-03-28 14:25:48,190 INFO codegen.CodeGenerator: Code generated in 14.860048 ms

Pickup_Borough	time_of_day	trip_count
EWR	night	3
EWR	afternoon	2
Unknown	morning	892
Unknown	afternoon	908
Unknown	evening	488
EWR	morning	5
Unknown	night	792

6.2 EVENING TRIP COUNTS BY PICKUP BOROUGH

Steps and APIs used:

1. *Filtering data:*
 - Reused the `borough_time_of_day_trip_counts` DataFrame created in 6.1.
 - Applied **`filter`** to select only rows where the "time_of_day" column value is "evening".
 - This ensures the focus on trips that occurred in the evening.
2. *Displaying results:*
 - Displayed the filtered Dataframe **`filtered_evening_trip_counts`** using **`show(truncate=False)`** to print dataframe without truncating column values, providing a clear view of pickup boroughs and their corresponding evening trip counts.

Knowledge gained:

- Learned how to filter DataFrames based on a specific value in a column.

Output:

2024-03-28 14:34:31,664 INFO codegen.CodeGenerator: Code generated in 15.684413 ms

Pickup_Borough	time_of_day	trip_count
Bronx	evening	1380355
Queens	evening	2223003
Manhattan	evening	5724796
Staten Island	evening	151276
Brooklyn	evening	3075616
Unknown	evening	488

6.3 TRIPS FROM BROOKLYN TO STATEN ISLAND**Steps and APIs used:**

1. *Data filtering:*
 - Performed **filter** on **nyc_df1** to select trips where the "Pickup_Borough" is "Brooklyn" and the "Dropoff_Borough" is "Staten Island".
 - This ensures the focus on trips that meet the specified origin and destination criteria.
2. *Selecting columns:*
 - Used **select** to choose only the "Pickup_Borough", "Dropoff_Borough", and "Pickup_Zone" columns for the results.
3. *Displaying and counting results:*
 - Used **show(10)** to display the first 10 rows of the filtered DataFrame **selected_df_final**.
 - Used **count** to determine the total number of trips that match the criteria and print the result using Python's **print** function.

Knowledge gained:

- Learned selecting specific columns using **select** for the output

Output:

2024-04-05 02:39:21,355 INFO scheduler.DAGScheduler: Job 4 finished: count at NativeMethodAccessorImpl.java:0, took 530.074293 s
→ Count of trips from Brooklyn to Staten Island: 69437
2024-04-05 02:39:21,387 INFO server.AbstractConnector: Stopped Spark@7574f7b1{HTTP/1.1,[http/1.1]}{0.0.0.0:4040}

2024-04-05 02:09:23,464 INFO codegen.CodeGenerator: Code generated in 15.135113 ms

Pickup_Borough	Dropoff_Borough	Pickup_Zone
Brooklyn	Staten Island	Columbia Street
Brooklyn	Staten Island	Columbia Street
Brooklyn	Staten Island	Columbia Street
Brooklyn	Staten Island	Columbia Street
Brooklyn	Staten Island	Columbia Street
Brooklyn	Staten Island	Marine Park/Mill ...
Brooklyn	Staten Island	Marine Park/Mill ...
Brooklyn	Staten Island	Marine Park/Mill ...
Brooklyn	Staten Island	Marine Park/Mill ...
Brooklyn	Staten Island	Marine Park/Mill ...

only showing top 10 rows

Task 7 Route Analysis

7.1 TOP 10 POPULAR ROUTES ANALYSIS

Steps and APIs used:

1. Data Loading and Preparation (Covered in Task 1 Report)

- Created `SparkSession`.
- Rideshare data and taxi zone lookup data are loaded from S3 using `textFile` and converted into `DataFrames`.
- Cleaned `taxi_df` using `regexp_replace`.
- `DataFrames` are joined based on pickup and dropoff locations using `join` and columns are renamed.
- The date column is formatted to "yyyy-MM-dd" format using `date_format`

2. Creating a new 'Route' column:

- Used `concat_ws` to create a new column using `withColumn`, named "Route" that combines the "Pickup_zone" and "Dropoff_zone" columns separated by "to".
- This creates a unique identifier for each route.

3. Calculating trip counts by route and business:

- `filter` to separate trips for Uber and Lyft businesses.
- Within each business group, used `groupBy` on "Route" and `count` to calculate the number of trips for each route.
- This results in two `DataFrames`: `route_profits_uber` (Uber trips) and `route_profits_lyft` (Lyft trips).

4. Merging DataFrames:

- Used `join` on the "Route" column to merge `route_profits_uber` and `route_profits_lyft` `DataFrames`.
- This creates a `DataFrame` `route_profits_merged` that includes "Route", "uber_count", and "lyft_count" columns.

5. Calculating total trip count:

- Created **total_count** column using **withColumn**. This column is a result of the addition of "uber_count", and "lyft_count" columns.
- Sorted the resulting Dataframe **route_profits_final** using **orderBy(by = "total_count", ascending = False)** in descending order on the basis of "total_count" column

6. Displaying results:

- Displayed the first 10 rows of route_profits_merged using **show(10, truncate=False)** to print without truncating column names, providing insights into the top 10 routes and their Uber/Lyft trip counts.

7. API's used:

- concat_ws()
- withColumn()
- filter()
- groupBy()
- agg()
- count()
- join()
- orderBy()

Knowledge gained:

- Learned how to create a new column by concatenating existing columns.
- It highlights grouping and aggregation using business type and route.
- Also learned how to merge DataFrames based on a shared column.

Output:

2024-04-05 03:22:41,916 INFO codegen.CodeGenerator: Code generated in 12.875848 ms

Route	uber_count	lyft_count	total_count
JFK Airport to NA	253211	46	253257
East New York to East New York	202719	184	202903
Borough Park to Borough Park	155803	78	155881
LaGuardia Airport to NA	151521	41	151562
Canarsie to Canarsie	126253	26	126279
South Ozone Park to JFK Airport	107392	1770	109162
Crown Heights North to Crown Heights North	98591	100	98691
Bay Ridge to Bay Ridge	98274	300	98574
Astoria to Astoria	90692	75	90767
Jackson Heights to Jackson Heights	89652	19	89671

only showing top 10 rows

Task 8 Graph Processing

8.1 DEFINING STRUCTYPE OF VERTEX AND EDGE SCHEMA

Steps and API's used:

1. StructType Definition (Spark SQL Structtype API):

- **StructType** is used to define the schema for both vertices and edges.
- Each schema is a list of StructField objects, specifying the **column name**, data type, and nullability.

2. Schema Definitions

- *Vertex Schema (vertex_Schema):*
 - *id (StringType, Not Null):* Unique identifier for each vertex (taxi zone). This field cannot be null to ensure every zone has a distinct ID.
 - *Borough (StringType, Nullable):* Borough name where the taxi zone is located (e.g., "Manhattan", "Queens"). This field can be null if the information is unavailable.
 - *Zone (StringType, Nullable):* Zone identifier within the borough. This field can be null if the information is unavailable.
 - *service_zone (StringType, Nullable):* Service zone identifier associated with the taxi zone. This field can be null if the information is unavailable.
- *Edge Schema (edge_Schema):*
 - *src (StringType, Not Null):* Source vertex ID (pickup location taxi zone). This field cannot be null to establish a starting point for the ride.
 - *dst (StringType, Not Null):* Destination vertex ID (dropoff location taxi zone). This field cannot be null to define the ride's endpoint.

Explanation:

1. The **vertexSchema** is established to define the structure of the vertices dataframe, which represents various locations or taxi zones.
2. The **edgeSchema** is defined to outline the structure of the edges dataframe, which captures the rideshare trips between locations.

Knowledge Gained:

- Learned defining schemas for vertices and edges with the Spark SQL **StructType API**.
- Learned how transform dataframes into a graph structure.

8.2 CONSTRUCT EDGES, VERTICES DATAFRAMES AND GRAPH

This section details the steps and APIs used to construct the vertices and edges DataFrames.

Steps and API's used:

1. *Vertices DataFrame:*

- Selected the relevant columns from the `taxi_df` DataFrame using the **select** function. These columns represent the properties of each location (vertex) in the graph.

2. *Edges DataFrame:*

- Selected the `pickup_location` and `dropoff_location` columns from the `ride_df` DataFrame using the **select** function. These columns define the connections (edges) between locations in the graph.

3. *GraphFrame Creation:*

- The selected vertices and edges DataFrames were used to construct a GraphFrame named **graph_main** using the **GraphFrame function**. This **GraphFrame** combines the vertex and edge information for **graph** analysis.

4. *Displaying Sample Data:*

- Utilized the **show** function on both **graph.vertices** and **graph.edges** to display the first 10 rows of each DataFrame.

5. *API's Used:*

- `select, alias`
- `GraphFrame`

Output:

```
2024-03-31 21:36:00,749 INFO scheduler.DAGScheduler: DAG finished, showing at nativeimage/iceberg-imp1.jar,v, took 0.00000 s
2024-03-31 21:36:00,749 INFO codegen.CodeGenerator: Code generated in 22.418454 ms
```

id	Borough	Zone	service_zone
1	EWB	Newark Airport	EWB
2	Queens	Jamaica Bay	Boro Zone
3	Bronx	Allerton/Pelham G...	Boro Zone
4	Manhattan	Alphabet City	Yellow Zone
5	Staten Island	Arden Heights	Boro Zone
6	Staten Island	Arrochar/Fort Wad...	Boro Zone
7	Queens	Astoria	Boro Zone
8	Queens	Astoria Park	Boro Zone
9	Queens	Auburndale	Boro Zone
10	Queens	Baisley Park	Boro Zone

only showing top 10 rows

Vertices Dataframe

```
2024-04-04 18:33:56,591 INFO codegen.CodeGenerator: Code generated in 21.367195 ms
```

src	dst
151	244
244	78
151	138
138	151
36	129
138	88
200	138
182	242
248	242
242	20

only showing top 10 rows

Edges Dataframe

8.3 CREATING A GRAPH USING VERTICES AND EDGES

Steps and API's used:

1. Finding Connected Vertices:

- Employed the **triplets** operation on the graph object with **distinct()** method. This operation transforms the GraphFrame into a DataFrame representation, where each row represents a connection (edge) in the graph and returns only unique rows.

2. Displaying Sample Data:

- Used the **show** function on the resulting DataFrame from the triplets operation using **graph.triplets.distinct().show(10, truncate=False)**. Kept the truncate parameter to False to ensure all data is displayed.

4. API's used:

- triplets
- show

Challenges:

- Finding right method to get output same as given.
- Got stuck with same similar rows, then resolved it by making use of distinct method.

Output:

```
2024-04-05 01:40:01,555 INFO scheduler.DAGScheduler: JOB # finished: showing at NativeMethodAccessorImpl.java:0, took 427.771675 s
2024-04-05 01:40:01,584 INFO codegen.CodeGenerator: Code generated in 12.22405 ms
```

src	edge	dst
[133, Brooklyn, Kensington, Boro Zone]	[133, 124]	[124, Queens, Howard Beach, Boro Zone]
[65, Brooklyn, Downtown Brooklyn/MetroTech, Boro Zone]	[65, 124]	[124, Queens, Howard Beach, Boro Zone]
[66, Brooklyn, DUMBO/Vinegar Hill, Boro Zone]	[66, 124]	[124, Queens, Howard Beach, Boro Zone]
[133, Brooklyn, Kensington, Boro Zone]	[133, 7]	[7, Queens, Astoria, Boro Zone]
[93, Queens, Flushing Meadows-Corona Park, Boro Zone]	[93, 7]	[7, Queens, Astoria, Boro Zone]
[34, Brooklyn, Brooklyn Navy Yard, Boro Zone]	[34, 234]	[234, Manhattan, Union Sq, Yellow Zone]
[256, Brooklyn, Williamsburg (South Side), Boro Zone]	[256, 234]	[234, Manhattan, Union Sq, Yellow Zone]
[223, Queens, Steinway, Boro Zone]	[223, 200]	[200, Bronx, Riverdale/North Riverdale/Fieldston, Boro Zone]
[47, Bronx, Claremont/Bathgate, Boro Zone]	[47, 200]	[200, Bronx, Riverdale/North Riverdale/Fieldston, Boro Zone]
[230, Manhattan, Times Sq/Theatre District, Yellow Zone]	[230, 200]	[200, Bronx, Riverdale/North Riverdale/Fieldston, Boro Zone]

only showing top 10 rows

8.4 COUNTING CONNECTED VERTICES WITH THE SAME BOROUGH AND SAME SERVICE ZONE

Steps and API's used:

1. *Extracting Graph Triples::*

- Utilized the **find** operation on the graph object to specify a pattern for finding connected vertices that is, **(a) - [e] -> (b)** which represents an edge **e** connecting vertex **a** to vertex **b**.

2. *Filtering by Borough and Service Zone:*

- Applied a **filter(a.Borough = b.Borough AND a.service_zone = b.service_zone)** on the result of the **find** operation. This filter expression ensures that only connections between vertices with the same borough and service zone are retained.

3. *Selecting and Renaming Columns:*

- Selected the desired columns (**a.id, b.id, a.Borough, and a.service_zone**) from the filtered DataFrame.

4. *Displaying Sample Data:*

- Used the **show** function on the resulting DataFrame from the select operation using **borough_service_vertices.distinct().show(10, truncate=False)**. Kept the truncate parameter to False to ensure all data is displayed and **distinct** method to get only unique values from **borough_service_vertices**.

5. *Counting and Displaying Results:*

- Used **count()** function on **borough_service_vertices** to get count of total connections and stored as **total_connected_vertices** including duplicates. Printed it out on terminal using print command as **print(f"Count of total connected vertices with the same Borough and service zone is: {total_connected_vertices}")**

4. API's used:

- find
- filter
- select
- distinct
- show

Knowledge gained:

- Learned how to traverse a graph by using the **find** operation to identify vertices connected by specific edge patterns.
- Use of **filter** to refine the search results and focus on specific pattern relevant to the task.
- Learned to remove duplicates using the **distinct** operation.

Output:

```
2024-04-04 19:10:06,225 INFO scheduler.DAGScheduler: Job 5 finished: showString at NativeMethodAccessorImpl.java:0, took 527.748352
2024-04-04 19:10:06,245 INFO codegen.CodeGenerator: Code generated in 11.262327 ms
```

```
+---+---+-----+-----+
|id |id |Borough  |service_zone|
+---+---+-----+-----+
|252|19 |Queens   |Boro Zone  |
|206|245|Staten Island|Boro Zone  |
|131|207|Queens   |Boro Zone  |
|111|178|Brooklyn |Boro Zone  |
|186|90 |Manhattan|Yellow Zone|
|64 |95 |Queens   |Boro Zone  |
|121|95 |Queens   |Boro Zone  |
|144|158|Manhattan|Yellow Zone|
|37 |65 |Brooklyn |Boro Zone  |
|164|68 |Manhattan|Yellow Zone|
+---+---+-----+-----+
```

only showing top 10 rows

```
2024-04-05 01:45:53,619 INFO scheduler.TaskSchedulerImpl: Killing all running tasks in stage 15: Stage finished
2024-04-05 01:45:53,619 INFO scheduler.DAGScheduler: Job 5 finished: count at NativeMethodAccessorImpl.java:0, took 351.519515 s
Count of total connected vertices with the same Borough and service zone is: 46886992
*****
2024-04-05 01:45:53,643 INFO server.AbstractConnector: Stopped Spark@9fe4561{HTTP/1.1,[http/1.1]}{0.0.0.0:4040}
2024-04-05 01:45:53,645 INFO ui.SparkUI: Stopped Spark web UI at http://task8-spark-app-40843c8eabe2119e-driver-svc.data-science-ec23
^
```

8.5 PERFORMING PAGE RANKING ON THE GRAPH DATAFRAME

Steps and API's used:

1. Applying PageRank:

- Utilized the **pageRank** function on the `graph_main` object. This function implements the PageRank algorithm, which iteratively calculates a score (PageRank) for each vertex, reflecting its relative importance based on the incoming links from other vertices

2. Setting Parameters:

- Set two parameters for the PageRank algorithm **resetProbability (0.17)** and **tol (0.01)** as instructed.
- resetProbability (0.17)**: This parameter represents the probability of a random jump to any vertex during the PageRank calculation. We used a value of 0.17.
- tol (0.01)**: This parameter defines the tolerance threshold for convergence. The algorithm iterates until the PageRank values for all vertices change by less than this threshold. We used a tolerance of 0.01.

3. Sorting and Displaying Results:

- Used the **sort** function on the result of the **pageRank** operation after selecting vertices.
- Sorted the vertices in descending order based on their calculated PageRank values (pagerank) using **sort('pagerank', ascending=False)** function.
- Then selected the columns `id` and `pagerank` using **select("id" , "pagerank")** and displayed the first 5 rows using **show** with `truncate=False`.

4. API's used:

- `pageRank`
- `sort`
- `select`
- `show`

Knowledge gained:

- This task demonstrated understanding of PageRank algorithm which is a valuable technique for ranking vertices in a graph based on their connectivity and influence within the network.

Output:

```
-----
2024-04-04 18:00:59,349 INFO scheduler.DAGScheduler: Job 40 finished: showString at NativeMethodAccessorImpl.java:0, toc
2024-04-04 18:00:59,366 INFO codegen.CodeGenerator: Code generated in 11.955213 ms
+---+-----+
|id |pagerank|
+---+-----+
|265|11.105433344107194|
|1  |5.4718454249167205|
|132|4.551132572067087 |
|138|3.5683223416564713|
|61 |2.6763973653412996|
+---+-----+
only showing top 5 rows
```